

Experiment Notebook

Setup Environment

```
In [1]: # DO NOT MODIFY THE CODE IN THIS CELL  
!pip install -q utstd  
  
from utstd.folders import *  
from utstd.ipyrenders import *  
  
at = AtFolder(  
    course_code=36106,  
    assignment="AT1",  
)  
at.run()  
  
import warnings  
warnings.simplefilter(action='ignore')
```

```
[notice] A new release of pip available: 22.3.1 -> 25.2  
[notice] To update, run: python.exe -m pip install --upgrade pip  
You can now save your data files in: c:\Users\brohao\Desktop\UTS\36106\AT1\36106\assignment\AT1\data
```

Student Information

```
In [2]: # Student name with space  
student_name = "Jiayu Hao"  
student_id = "25948860"
```

```
In [3]: # DO NOT MODIFY THE CODE IN THIS CELL  
print_tile(size="h1", key='student_name', value=student_name)  
  
student_name
```

Jiayu Hao

```
In [4]: # DO NOT MODIFY THE CODE IN THIS CELL  
print_tile(size="h1", key='student_id', value=student_id)  
  
student_id
```

25948860

0. Python Packages

0.a Install Additional Packages

If you are using additional packages, you need to install them here using the command: ! pip install <package_name>

```
In [5]: !pip install numpy  
!pip install matplotlib  
!pip install seaborn
```

```
Requirement already satisfied: numpy in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (2.3.2)  
[notice] A new release of pip available: 22.3.1 -> 25.2  
[notice] To update, run: python.exe -m pip install --upgrade pip  
ERROR: Could not find a version that satisfies the requirement matplotlib.lib (from versions: none)  
ERROR: No matching distribution found for matplotlib.lib  
  
[notice] A new release of pip available: 22.3.1 -> 25.2  
[notice] To update, run: python.exe -m pip install --upgrade pip  
Requirement already satisfied: seaborn in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (0.13.2)  
Requirement already satisfied: numpy!=1.24.0,>=1.20 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from seaborn) (2.3.2)  
Requirement already satisfied: pandas>=1.2 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from seaborn) (2.2.2)  
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from seaborn) (3.10.5)  
Requirement already satisfied: contourpy>=1.0.1 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.3)  
Requirement already satisfied: cycler>=0.10 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.59.1)  
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.9)  
Requirement already satisfied: packaging>=20.0 in c:\users\brohao\appdata\roaming\python\python311\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (25.0)  
Requirement already satisfied: pillow>=8 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (11.3.0)  
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.2.3)  
Requirement already satisfied: python-dateutil>=2.7 in c:\users\brohao\appdata\roaming\python\python311\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from pandas>=1.2->seaborn) (2023.3)  
Requirement already satisfied: six>=1.5 in c:\users\brohao\appdata\local\programs\python\python311\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)  
[notice] A new release of pip available: 22.3.1 -> 25.2  
[notice] To update, run: python.exe -m pip install --upgrade pip
```

0.b Import Packages

```
In [6]: # DO NOT MODIFY THE CODE IN THIS CELL  
import pandas as pd  
import altair as alt
```

```
In [7]: import matplotlib.pyplot as plt  
import seaborn as sns  
import numpy as np
```

A. Project Description

```
In [8]: # Explain clearly what is the goal of this project for the business. How will the results be used?  
  
business_objective = """  
The goal of this project is to predict how much customers will pay for their Motor Vehicle Insurance.  
The results will help the company set fair premiums, keep customers, and lower financial risk.  
If predictions are accurate, the company can give fair prices, stay profitable, and keep customers.  
"""
```

```
In [9]: # DO NOT MODIFY THE CODE IN THIS CELL  
print_tile(size="h3", key='business_objective', value=business_objective)
```

business_objective

The goal of this project is to predict how much customers will pay for their Motor Vehicle Insurance. The results will help the company set fair premiums, keep customers, and lower financial risk. If predictions are accurate, the company can give fair prices, stay profitable, and keep customers. If predictions are wrong, the company may lose money by charging too little or lose customers by charging too much.

B. Dataset Understanding (Global Interpretation)

```
In [10]: # DO NOT MODIFY THE CODE IN THIS CELL  
# Load training data  
try:  
    training_df = pd.read_csv(at.folder_path / "car_insurance_premium_training.csv")  
    validation_df = pd.read_csv(at.folder_path / "car_insurance_premium_validation.csv")  
    testing_df = pd.read_csv(at.folder_path / "car_insurance_premium_testing.csv")  
except Exception as e:  
    print(e)
```

B.1 Explore Training Set

You can add more cells in this section

```
In [11]: print("Shape of dataset:", training_df.shape)    # rows * columns  
# Overview of data types and non-null counts  
print("\nDataset Info:")  
print(training_df.info())  
  
# First few rows  
print("\nFirst 5 rows of training set:")  
pd.set_option("display.max_columns", None)  
display(training_df.head())  
  
# Summary statistics for numerical features  
print("\nSummary statistics for numerical variables:")  
display(training_df.describe().T)  
  
# Summary for categorical features  
print("\nUnique values count for categorical variables:")  
categorical_cols = training_df.select_dtypes(include=["object", "category"]).columns  
for col in categorical_cols:  
    print(f"{col}: {training_df[col].nunique()} unique values")
```

```

# Step 8: Missing values
print("\nMissing values per column:")
missing = training_df.isnull().sum()
missing = missing[missing > 0].sort_values(ascending=False)
display(missing)

# Step 9: Check for duplicate rows
print("\nNumber of duplicate rows:", training_df.duplicated().sum())

```

Shape of dataset: (32136, 40)

Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32136 entries, 0 to 32135
Data columns (total 40 columns):

```

#	Column	Non-Null Count	Dtype
0	customer_id	32136	object
1	prefix	21345	object
2	first_name	32136	object
3	last_name	32116	object
4	gender	32136	object
5	birth_date	32136	object
6	driving_license_date	32136	object
7	phone_number	32136	object
8	email	32136	object
9	secondary_address	32136	object
10	building_number	32136	int64
11	street_name	32136	object
12	street_suffix	32136	object
13	suburb	32136	object
14	contract_start_date	32136	object
15	last_renewal_date	32136	object
16	next_renewal_date	32136	object
17	distribution_channel	32136	object
18	seniority	32136	int64
19	current_policies_held	32136	int64
20	max_policies_held	32136	int64
21	max_products_held	32136	int64
22	lapsed_policies	32136	int64
23	lapsed_date	12567	object
24	payment_method	32136	int64
25	net_premium_amount	32136	float64
26	total_claims_cost_in_current_year	32136	float64
27	total_claims_number_in_current_year	32136	int64
28	total_claims_number_in_history	32136	int64
29	total_claims_number_ratio	32136	float64
30	policy_type	32136	int64
31	second_driver	32136	int64
32	matriculation_year	32136	int64
33	vehicle_horsepower	32136	int64
34	vehicle_cylinder	32136	int64
35	vehicle_value	32136	float64
36	vehicle_doors	32136	int64
37	vehicle_fuel_type	32058	object
38	vehicle_length	31451	float64
39	vehicle_weight	32136	int64

dtypes: float64(5), int64(16), object(19)

memory usage: 9.8+ MB

None

First 5 rows of training set:

	customer_id	prefix	first_name	last_name	gender	birth_date	driving_license_date	phone_n
0	2f4cea69-3806-41b7-b7c2-6039eaac8fae	Mr.	John	Medina	m	1970-12-24	2004-12-21	+61 4
1	c738e8cb-13fe-43f9-8d59-b501b04ff590	Dr.	Krystal	Howard	f	1989-02-07	2008-02-20	0412-9
2	685c81d7-b1de-4862-902e-f57612dedb57	Mr.	Jeffrey	Craig	m	1972-08-29	2001-12-08	0790
3	643be7d5-6ffa-4517-8468-85bfe5e01bde	Mr.	Gregory	Ortiz	m	1983-01-01	2013-09-25	(03) 963
4	99c498eb-33de-479b-88ec-b5b9d6d5252b	Mr.	Donald	Harrison	m	1960-07-16	1999-06-30	+61-2

◀  ▶ Summary statistics for numerical variables:

		count	mean	std	min	25%
	building_number	32136.0	185.629792	280.506483	0.00	6.000
	seniority	32136.0	5.896596	5.909002	1.00	2.000
	current_policies_held	32136.0	1.443055	0.943590	1.00	1.000
	max_policies_held	32136.0	1.796023	1.157598	1.00	1.000
	max_products_held	32136.0	1.048015	0.231683	1.00	1.000
	lapsed_policies	32136.0	0.391804	0.550027	0.00	0.000
	payment_method	32136.0	0.361433	0.480423	0.00	0.000
	net_premium_amount	32136.0	579.367438	44.991023	460.00	542.000
	total_claims_cost_in_current_year	32136.0	104.826855	1110.840111	0.00	0.000
	total_claims_number_in_current_year	32136.0	0.206622	0.771437	0.00	0.000
	total_claims_number_in_history	32136.0	2.647498	3.787398	0.00	0.000
	total_claims_number_ratio	32136.0	0.463897	0.774516	0.00	0.000
	policy_type	32136.0	2.839806	0.396159	1.00	3.000
	second_driver	32136.0	0.119461	0.324335	0.00	0.000
	matriculation_year	32136.0	2004.573376	5.347809	1954.00	2001.000
	vehicle_horsepower	32136.0	97.543409	28.174377	0.00	75.000
	vehicle_cylinder	32136.0	1717.529842	409.694777	49.00	1461.000
	vehicle_value	32136.0	18972.391870	7176.732691	270.46	14093.730
	vehicle_doors	32136.0	4.404998	0.976780	0.00	4.000
	vehicle_length	31451.0	4.251559	0.378661	2.50	4.015
	vehicle_weight	32136.0	1252.991349	274.112493	55.00	1089.000

◀ Unique values count for categorical variables:

customer_id: 32136 unique values
prefix: 5 unique values
first_name: 1321 unique values
last_name: 2446 unique values
gender: 3 unique values
birth_date: 14183 unique values
driving_license_date: 11691 unique values
phone_number: 32136 unique values
email: 30167 unique values
secondary_address: 3270 unique values
street_name: 27443 unique values
street_suffix: 200 unique values
suburb: 18887 unique values
contract_start_date: 4963 unique values
last_renewal_date: 1117 unique values
next_renewal_date: 1117 unique values
distribution_channel: 3 unique values
lapsed_date: 1200 unique values
vehicle_fuel_type: 2 unique values

Missing values per column:

```
lapsed_date      19569
prefix           10791
vehicle_length   685
vehicle_fuel_type 78
last_name        20
dtype: int64
Number of duplicate rows: 0
```

```
In [12]: training_set_insights = """
The dataset has 32,136 rows and 40 columns, including 3 types information(customer, policy, and vehicle information). Several columns are identifiers or personal details(IDs, names, contact info, addresses). Those columns will be dropped. Missing values exist mainly in lapsed_date(61%), prefix(34%), and a few vehicle information(vehicle_length, vehicle_fuel_type).
Numerical issues:
Outliers: total_claims_cost_in_current_year(mean: 105, max:128, 000).
Invalid values (e.g., vehicle_doors=0, vehicle_weight=55, vehicle_horsepower=0).
Skewed variables (claims cost(exist outlier), vehicle_value).

Categorical variables: gender, policy_type, vehicle_fuel_type, and distribution_channel will be encoded. Some columns may be not useful. Dates(birth_date, driving_license_date, matriculation_year) will be transformed into derived features such as age, driving experience, and car age.
"""

In [13]: # DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='training_set_insights', value=training_set_insights)
```

training_set_insights

The dataset has 32,136 rows and 40 columns, including 3 types information(customer, policy, and vehicle information). Several columns are identifiers or personal details(IDs, names, contact info, addresses). Those columns will be dropped. Missing values exist mainly in lapsed_date(61%), prefix(34%), and a few vehicle information(vehicle_length, vehicle_fuel_type). Numerical issues: Outliers: total_claims_cost_in_current_year(mean: 105, max:128, 000). Invalid values (e.g., vehicle_doors=0, vehicle_weight=55, vehicle_horsepower=0). Skewed variables (claims cost(exist outlier), vehicle_value). Categorical variables: gender, policy_type, vehicle_fuel_type, and distribution_channel will be encoded. Some columns may be not useful. Dates(birth_date, driving_license_date, matriculation_year) will be transformed into derived features such as age, driving experience, and car age.

B.2 Explore Validation Set

You can add more cells in this section

```
In [14]: print("Shape of dataset:", validation_df.shape)    # rows * columns

# Overview of data types and non-null counts
print("\nDataset Info:")
print(validation_df.info())

# First few rows
print("\nFirst 5 rows of validation set:")
pd.set_option("display.max_columns", None)
display(validation_df.head())

# Summary statistics for numerical features
print("\nSummary statistics for numerical variables:")
display(validation_df.describe().T)
```

```
# Summary for categorical features
print("\nUnique values count for categorical variables:")
categorical_cols = validation_df.select_dtypes(include=["object", "category"]).columns
for col in categorical_cols:
    print(f"{col}: {validation_df[col].nunique()} unique values")

# Missing values
print("\nMissing values per column:")
missing = validation_df.isnull().sum()
missing = missing[missing > 0].sort_values(ascending=False)
display(missing)

# Duplicate rows
print("\nNumber of duplicate rows:", validation_df.duplicated().sum())
```

Shape of dataset: (10700, 40)

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10700 entries, 0 to 10699

Data columns (total 40 columns):

#	Column	Non-Null Count	Dtype
0	customer_id	10700	non-null object
1	prefix	7078	non-null object
2	first_name	10700	non-null object
3	last_name	10696	non-null object
4	gender	10700	non-null object
5	birth_date	10700	non-null object
6	driving_license_date	10700	non-null object
7	phone_number	10700	non-null object
8	email	10700	non-null object
9	secondary_address	10700	non-null object
10	building_number	10700	non-null int64
11	street_name	10700	non-null object
12	street_suffix	10700	non-null object
13	suburb	10700	non-null object
14	contract_start_date	10700	non-null object
15	last_renewal_date	10700	non-null object
16	next_renewal_date	10700	non-null object
17	distribution_channel	10700	non-null object
18	seniority	10700	non-null int64
19	current_policies_held	10700	non-null int64
20	max_policies_held	10700	non-null int64
21	max_products_held	10700	non-null int64
22	lapsed_policies	10700	non-null int64
23	lapsed_date	3087	non-null object
24	payment_method	10700	non-null int64
25	net_premium_amount	10700	non-null float64
26	total_claims_cost_in_current_year	10700	non-null float64
27	total_claims_number_in_current_year	10700	non-null int64
28	total_claims_number_in_history	10700	non-null int64
29	total_claims_number_ratio	10700	non-null float64
30	policy_type	10700	non-null int64
31	second_driver	10700	non-null int64
32	matriculation_year	10700	non-null int64
33	vehicle_horsepower	10700	non-null int64
34	vehicle_cylinder	10700	non-null int64
35	vehicle_value	10700	non-null float64
36	vehicle_doors	10700	non-null int64
37	vehicle_fuel_type	10081	non-null object
38	vehicle_length	7261	non-null float64
39	vehicle_weight	10700	non-null int64

dtypes: float64(5), int64(16), object(19)

memory usage: 3.3+ MB

None

First 5 rows of validation set:

	customer_id	prefix	first_name	last_name	gender	birth_date	driving_license_date	phone_n
0	980657d0-843c-4006-abcb-aa6a44ad57f1	Nan	Veronica	Simpson	u	1982-08-12	2000-09-21	560
1	f50db3af-dfae-496a-8aaf-14b0ea2729e2	Dr.	Carla	Riley	f	1976-12-09	1994-12-23	+61 4
2	d32a95a8-3ed8-4a98-a54a-472dd480f6fa	Nan	Brian	Branch	u	1948-06-25	1970-04-11	(07).963
3	ba0d15a9-0792-44a2-87d3-05f45c1c046d	Dr.	Drew	Smith	m	1974-09-21	1994-09-21	02.498
4	e0da2a4b-fa0f-4512-a363-780d1f8e2b0e	Mr.	Jose	Watson	m	1961-06-26	2014-08-01	(07).751

◀  ▶ Summary statistics for numerical variables:

		count	mean	std	min	25%
	building_number	10700.0	185.951776	280.713477	0.000	6.000
	seniority	10700.0	6.120374	6.444548	1.000	2.000
	current_policies_held	10700.0	1.756449	1.268339	1.000	1.000
	max_policies_held	10700.0	2.082991	1.410385	1.000	1.000
	max_products_held	10700.0	1.108411	0.336329	1.000	1.000
	lapsed_policies	10700.0	0.315327	0.558235	0.000	0.000
	payment_method	10700.0	0.161682	0.368176	0.000	0.000
	net_premium_amount	10700.0	489.297009	133.399129	309.000	396.000
	total_claims_cost_in_current_year	10700.0	87.516426	2356.273613	0.000	0.000
	total_claims_number_in_current_year	10700.0	0.156822	0.659318	0.000	0.000
	total_claims_number_in_history	10700.0	2.149533	3.518325	0.000	0.000
	total_claims_number_ratio	10700.0	0.322316	0.816566	0.000	0.000
	policy_type	10700.0	2.387383	0.929685	1.000	1.000
	second_driver	10700.0	0.084766	0.278547	0.000	0.000
	matriculation_year	10700.0	2002.933271	9.513792	1950.000	2000.000
	vehicle_horsepower	10700.0	70.655421	44.015785	0.000	34.000
	vehicle_cylinder	10700.0	1296.187477	867.199458	49.000	599.000
	vehicle_value	10700.0	14177.850521	10079.413937	270.460	6611.130
	vehicle_doors	10700.0	3.056822	2.178017	0.000	0.000
	vehicle_length	7261.0	4.151021	0.374925	1.978	3.916
	vehicle_weight	10700.0	956.078411	710.277947	43.000	200.000

◀ Unique values count for categorical variables:

customer_id: 10700 unique values
prefix: 5 unique values
first_name: 1017 unique values
last_name: 1624 unique values
gender: 3 unique values
birth_date: 7528 unique values
driving_license_date: 6625 unique values
phone_number: 10700 unique values
email: 10445 unique values
secondary_address: 2516 unique values
street_name: 10139 unique values
street_suffix: 200 unique values
suburb: 8240 unique values
contract_start_date: 3116 unique values
last_renewal_date: 1056 unique values
next_renewal_date: 1056 unique values
distribution_channel: 3 unique values
lapsed_date: 1027 unique values
vehicle_fuel_type: 2 unique values

Missing values per column:

```
lapsed_date      7613
prefix          3622
vehicle_length   3439
vehicle_fuel_type 619
last_name         4
dtype: int64
Number of duplicate rows: 0
```

```
In [15]: # provide a detailed analysis on the validation set, its dimensions, information, issues, ...
validation_set_insights = """
The validation set has 10,700 rows and 40 columns, same structure as training.
Premiums are lower on average (mean ≈ 489) and show more variation.
Some variables have many missing values (lapsed_date, prefix, vehicle_length).
several contain outliers (e.g., very high claims, unusual vehicle specs).
Data quality issues are similar to training and need the same cleaning.
Overall, the set is consistent and suitable for model evaluation.
"""

In [16]: # DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='validation_set_insights', value=validation_set_insights)
```

validation_set_insights

The validation set has 10,700 rows and 40 columns, same structure as training. Premiums are lower on average (mean \approx 489) and show more variation. Some variables have many missing values (lapsed_date, prefix, vehicle_length). several contain outliers (e.g., very high claims, unusual vehicle specs). Data quality issues are similar to training and need the same cleaning. Overall, the set is consistent and suitable for model evaluation.

B.3 Explore Testing Set

You can add more cells in this section

```
In [17]: print("Shape of dataset:", testing_df.shape)    # rows * columns

# Overview of data types and non-null counts
print("\nDataset Info:")
print(testing_df.info())

# First few rows
print("\nFirst 5 rows of testing set:")
pd.set_option("display.max_columns", None)
display(testing_df.head())

# Summary statistics for numerical features
print("\nSummary statistics for numerical variables:")
display(testing_df.describe().T)

# Summary for categorical features
print("\nUnique values count for categorical variables:")
categorical_cols = testing_df.select_dtypes(include=["object", "category"]).columns
for col in categorical_cols:
    print(f"{col}: {testing_df[col].nunique()} unique values")

# Missing values
print("\nMissing values per column:")
missing = testing_df.isnull().sum()
missing = missing[missing > 0].sort_values(ascending=False)
display(missing)
```

```
# Duplicate rows
print("\nNumber of duplicate rows:", testing_df.duplicated().sum())
```

Shape of dataset: (10666, 40)

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10666 entries, 0 to 10665
Data columns (total 40 columns):
```

#	Column	Non-Null Count	Dtype
0	customer_id	10666	object
1	prefix	7039	object
2	first_name	10666	object
3	last_name	10665	object
4	gender	10666	object
5	birth_date	10666	object
6	driving_license_date	10666	object
7	phone_number	10666	object
8	email	10666	object
9	secondary_address	10666	object
10	building_number	10666	int64
11	street_name	10666	object
12	street_suffix	10666	object
13	suburb	10666	object
14	contract_start_date	10666	object
15	last_renewal_date	10666	object
16	next_renewal_date	10666	object
17	distribution_channel	10666	object
18	seniority	10666	int64
19	current_policies_held	10666	int64
20	max_policies_held	10666	int64
21	max_products_held	10666	int64
22	lapsed_policies	10666	int64
23	lapsed_date	4539	object
24	payment_method	10666	int64
25	net_premium_amount	10666	float64
26	total_claims_cost_in_current_year	10666	float64
27	total_claims_number_in_current_year	10666	int64
28	total_claims_number_in_history	10666	int64
29	total_claims_number_ratio	10666	float64
30	policy_type	10666	int64
31	second_driver	10666	int64
32	matriculation_year	10666	int64
33	vehicle_horsepower	10666	int64
34	vehicle_cylinder	10666	int64
35	vehicle_value	10666	float64
36	vehicle_doors	10666	int64
37	vehicle_fuel_type	10506	object
38	vehicle_length	9648	float64
39	vehicle_weight	10666	int64

dtypes: float64(5), int64(16), object(19)

memory usage: 3.3+ MB

None

First 5 rows of testing set:

	customer_id	prefix	first_name	last_name	gender	birth_date	driving_license_date	phone_r
0	08630f39-3b19-4c50-aaac-6ccb5d63b820	Mr.	Jose	Fritz	m	1998-07-15	2016-11-22	+61.7.521
1	e491ab7e-a3df-474d-a6d5-9cbd85dc4242	NaN	Eric	Wilson	u	1992-05-12	2010-11-05	812
2	6bafa529-bdc5-442f-b96f-eb5120799054	Mr.	Parker	Burton	m	1957-04-07	1977-04-24	0463..
3	aea650ed-62e6-4b69-99b3-04fe84f60354	NaN	Matthew	Bonilla	u	1981-04-14	2008-03-14	+61.477..
4	fb24f594-dc26-4cb5-8fdc-c8568b3c72ea	Mrs.	Sarah	Holloway	f	1974-08-01	1992-09-29	0411..

◀  ▶ Summary statistics for numerical variables:

		count	mean	std	min	25%
	building_number	10666.0	186.816801	281.124963	0.00	6.000
	seniority	10666.0	5.153666	5.258575	1.00	2.000
	current_policies_held	10666.0	1.483218	0.966358	1.00	1.000
	max_policies_held	10666.0	1.828146	1.166254	1.00	1.000
	max_products_held	10666.0	1.056722	0.245481	1.00	1.000
	lapsed_policies	10666.0	0.429402	0.565399	0.00	0.000
	payment_method	10666.0	0.450778	0.497595	0.00	0.000
	net_premium_amount	10666.0	773.221733	212.915798	397.00	694.000
	total_claims_cost_in_current_year	10666.0	184.485467	1288.970159	0.00	0.000
	total_claims_number_in_current_year	10666.0	0.267110	0.899549	0.00	0.000
	total_claims_number_in_history	10666.0	2.952466	4.189095	0.00	0.000
	total_claims_number_ratio	10666.0	0.618153	0.929463	0.00	0.000
	policy_type	10666.0	2.705888	0.619905	1.00	3.000
	second_driver	10666.0	0.194825	0.396084	0.00	0.000
	matriculation_year	10666.0	2007.177386	6.766693	1952.00	2003.000
	vehicle_horsepower	10666.0	103.245265	44.324701	0.00	82.000
	vehicle_cylinder	10666.0	1645.733827	627.110106	49.00	1396.000
	vehicle_value	10666.0	21701.601107	11886.970952	480.81	15308.335
	vehicle_doors	10666.0	4.090381	1.527644	0.00	4.000
	vehicle_length	9648.0	4.356239	0.429827	2.50	4.110
	vehicle_weight	10666.0	1238.638009	435.880734	43.00	1095.000

◀ Unique values count for categorical variables:

customer_id: 10666 unique values
prefix: 5 unique values
first_name: 1025 unique values
last_name: 1664 unique values
gender: 3 unique values
birth_date: 7753 unique values
driving_license_date: 6827 unique values
phone_number: 10666 unique values
email: 10448 unique values
secondary_address: 2542 unique values
street_name: 10127 unique values
street_suffix: 200 unique values
suburb: 8226 unique values
contract_start_date: 2805 unique values
last_renewal_date: 1095 unique values
next_renewal_date: 1095 unique values
distribution_channel: 3 unique values
lapsed_date: 1111 unique values
vehicle_fuel_type: 2 unique values

Missing values per column:

```
lapsed_date      6127
prefix          3627
vehicle_length   1018
vehicle_fuel_type 160
last_name         1
dtype: int64
Number of duplicate rows: 0
```

```
In [18]: # provide a detailed analysis on the testing set, its dimensions, information, issues, ...
testing_set_insights = """
The testing set has 10,666 rows and 40 columns, consistent with training and validation.
Average premium is higher (mean ≈ 773) and shows more variability compared to other sets.
Several variables contain missing values, most notably lapsed_date (~57%), prefix (~34%), and
Outliers exist in claims (e.g. very high costs), vehicle specs (e.g. horsepower up to 580, value 200k),
Despite these issues, the structure matches the other datasets, making it suitable for final performance evaluation after consistent preprocessing.
"""

In [19]: # DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='testing_set_insights', value=testing_set_insights)
```

testing_set_insights

The testing set has 10,666 rows and 40 columns, consistent with training and validation. Average premium is higher (mean \approx 773) and shows more variability compared to other sets. Several variables contain missing values, most notably lapsed_date (~57%), prefix (~34%), and vehicle_length (~10%). Outliers exist in claims (e.g. very high costs), vehicle specs (e.g. horsepower up to 580, value 200k), and unrealistic cases (vehicle_doors=0). Despite these issues, the structure matches the other datasets, making it suitable for final performance evaluation after consistent preprocessing.

C. Feature Understanding (Local Interpretation)

C.1 Explore Target Variable

Save the name of column used as the target variable and call it `target_name`

You can add more cells in this section

```
In [20]: # Target name
target_name = 'net_premium_amount'
```

```
In [21]: # Basic statistics
target = training_df[target_name]
print("Target Variable: net_premium_amount")
print(target.describe())

# Histogram + KDE
plt.figure(figsize=(10,5))
sns.histplot(target, bins=30, kde=True, color="steelblue")
plt.title("Distribution of net_premium_amount", fontsize=14)
plt.xlabel("Net Premium Amount")
plt.ylabel("Frequency")
plt.show()

# Boxplot (to check for outliers)
plt.figure(figsize=(8,3))
```


```
sns.boxplot(x=target, color="lightcoral")
plt.title("Boxplot of Net Premium Amount", fontsize=14)
plt.xlabel("Net Premium Amount")
plt.show()
```

Target Variable: net_premium_amount


count 32136.000000
mean 579.367438
std 44.991023
min 460.000000
25% 542.000000
50% 575.000000
75% 614.000000
max 676.000000

Name: net_premium_amount, dtype: float64

Distribution of net_premium_amount



Boxplot of Net Premium Amount



```
In [22]: # provide a detailed analysis on the target variable, its distribution, limitations, issues,
target_insights = """
The target variable net_premium_amount ranges from 460 to 676, with a mean of about 579 and a
Its distribution is relatively narrow and right-skewed, most values of target variable concen
This limited variability means even small errors (e.g., ±20) are significant compared to the p
No extreme outliers are observed, means the premium scale need to be limited strictly.
"""
```

```
In [23]: # DO NOT MODIFY THE CODE IN THIS CELL  
print_tile(size="h3", key='target_insights', value=target_insights)
```

target_insights

The target variable net_premium_amount ranges from 460 to 676, with a mean of about 579 and a small standard deviation (~45). Its distribution is relatively narrow and right-skewed, most values of target variable concentrate between 540 and 615. This limited variability means even small errors (e.g., ± 20) are significant compared to the premium scale. No extreme outliers are observed, means the premium scale need to be limited strictly.

C.2 Explore Feature of Interest birth_date

You can add more cells in this section

```
In [24]: feature_name = "birth_date"          # Birthday  
ref_date_col = "contract_start_date"    # Contract start date  
target_name = "net_premium_amount"       # Target variable  
print_feature_name = "age (at contract start)"
```

```
In [25]: # Parse datetimes  
birth = pd.to_datetime(training_df[feature_name], errors="coerce")  
contract_start = pd.to_datetime(training_df[ref_date_col], errors="coerce")  
  
# Compute age (years) at contract start  
year_diff = contract_start.dt.year - birth.dt.year  
before_birthday = (  
    (contract_start.dt.month < birth.dt.month) |  
    ((contract_start.dt.month == birth.dt.month) & (contract_start.dt.day < birth.dt.day))  
)  
age_at_contract = year_diff - before_birthday.astype(int)  
age_at_contract = age_at_contract.where(birth.notna() & contract_start.notna(), np.nan)  
  
# Basic stat  
print(f"Feature: {feature_name} -> {print_feature_name}")  
print(age_at_contract.describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]))  
  
# Reasonableness checks  
if (age_at_contract < 0).any():  
    print("Warning: Negative ages found.")  
if (age_at_contract < 16).any():  
    print("Warning: Ages < 16 detected.")  
if (age_at_contract > 100).any():  
    print("Warning: Ages > 100 detected.")  
  
# Distribution plots  
valid_age = age_at_contract.dropna()  
plt.figure(figsize=(10,5))  
plt.hist(valid_age, bins=30)  
plt.title(f"{print_feature_name} Distribution")  
plt.xlabel(f"{print_feature_name}")  
plt.ylabel("Count")  
plt.show()  
  
plt.figure(figsize=(8,3))  
plt.boxplot(valid_age, vert=False)  
plt.title(f"Boxplot of {print_feature_name}")  
plt.xlabel(f"{print_feature_name}")  
plt.show()
```

```

# Relationship with target
df_plot = pd.DataFrame({"Age_at_Contract": age_at_contract, target_name: training_df[target_name]})

# Scatter plot
plt.figure(figsize=(8,5))
plt.scatter(df_plot["Age_at_Contract"], df_plot[target_name], alpha=0.3)
plt.title(f"{print_feature_name} vs Net Premium Amount")
plt.xlabel(f"{print_feature_name}")
plt.ylabel("Net Premium Amount")
plt.show()


# Grouped averages
df_plot["Age_group"] = pd.cut(df_plot["Age_at_Contract"], bins=[0,25,35,50,65,100], right=False)
mean_prem = df_plot.groupby("Age_group")[target_name].mean()

plt.figure(figsize=(7,4))
plt.bar(mean_prem.index.astype(str), mean_prem.values)
plt.title(f"Mean Premium by {print_feature_name}")
plt.ylabel("Mean Premium"); plt.xlabel("Age Group")
plt.xticks(rotation=45, ha="right")
plt.show()


```

Feature: birth_date -> age (at contract start)


count	32136.000000
mean	43.933408
std	12.497067
min	18.000000
1%	20.350000
5%	26.000000
25%	34.000000
50%	43.000000
75%	53.000000
95%	66.000000
99%	73.000000
max	91.000000
dtype:	float64




Boxplot of age (at contract start)



age (at contract start) vs Net Premium Amount



Mean Premium by age (at contract start)



```
In [26]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues,
feature_1_insights = """
Age distribution from 18-elder(90+), right-skewed toward middle-aged drivers.
Younger drivers are usually higher risk -> higher premiums; older drivers may have mixed risk
Issues: the original data 'birth_date'(string) must be converted to age(at contract start).
It's continuous variable.
"""
```

```
In [27]: # DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='feature_1_insights', value=feature_1_insights)
```

feature_1_insights

Age distribution from 18-elder(90+), right-skewed toward middle-aged drivers. Younger drivers are usually higher risk - higher premiums; older drivers may have mixed risk depending on some other features. Issues: the original data 'birth_date'(string) must be converted to age(at contract start). It's continuous variable.

C.3 Explore Feature of Interest driving_license_date

You can add more cells in this section

```
In [28]: feature_name      = "driving_license_date"          # License issue date
ref_date_col       = "contract_start_date"           # Reference date (contract start)
target_name        = "net_premium_amount"            # Target variable
print_feature_name = "Driving Experience (years) at Contract Start"
```

```
In [29]: # Parse datetimes
lic = pd.to_datetime(training_df[feature_name], errors="coerce")
ref = pd.to_datetime(training_df[ref_date_col], errors="coerce")

# Compute driving experience (years) at contract start (allow negatives)
year_diff = ref.dt.year - lic.dt.year
before_anniv = (
```

```

        (ref.dt.month < lic.dt.month) |
        ((ref.dt.month == lic.dt.month) & (ref.dt.day < lic.dt.day))
    )
exp_years = (year_diff - before_anniv.astype(int)).where(lic.notna() & ref.notna(), np.nan)

# Descriptive statistics
print(f"Feature: {feature_name} -> {print_feature_name}")
print(exp_years.describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]))


# Distribution plots (include negatives)
valid = exp_years.dropna()
if len(valid) > 0:
    plt.figure(figsize=(10,5))
    plt.hist(valid, bins=30)
    plt.title(print_feature_name)
    plt.xlabel("Years"); plt.ylabel("Count")
    plt.axvline(0, linestyle="--") # reference line at 0 year (optional)
    plt.tight_layout(); plt.show()

    plt.figure(figsize=(8,3))
    plt.boxplot(valid, vert=False)
    plt.title(f"Boxplot: {print_feature_name}")
    plt.xlabel("Years")
    plt.tight_layout(); plt.show()


# Relationship with target
df_plot = pd.DataFrame({"Experience": exp_years, target_name: training_df[target_name]}).dropna()

# Scatter
plt.figure(figsize=(8,5))
plt.scatter(df_plot["Experience"], df_plot[target_name], alpha=0.3)
plt.title(f"{print_feature_name} vs {target_name}")
plt.xlabel("Experience (years)"); plt.ylabel(target_name)
plt.tight_layout(); plt.show()



# Grouped averages (bins include negatives)
# Adjust bins as needed for your data distribution
lo, hi = np.nanpercentile(df_plot["Experience"], [1, 99])
bins = [-50, -1, 0, 2, 5, 10, 20, 40, max(80, hi)]
df_plot["ExpGroup"] = pd.cut(df_plot["Experience"], bins=bins, right=False)
mean_prem = df_plot.groupby("ExpGroup")[target_name].mean()


plt.figure(figsize=(9,4))
plt.bar(mean_prem.index.astype(str), mean_prem.values)
plt.title(f"Mean {target_name} by {print_feature_name} Group")
plt.ylabel(f"Mean {target_name}"); plt.xlabel("Experience Group (years)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout(); plt.show()

```

Feature: driving_license_date -> Driving Experience (years) at Contract Start

count	32136.000000
mean	21.090864
std	12.114187
min	-13.000000
1%	0.000000
5%	4.000000
25%	11.000000
50%	20.000000
75%	30.000000
95%	43.000000
99%	50.000000
max	71.000000
dtype:	float64





```
In [30]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues, etc.

feature_2_insights = """
Driving Experience (years) at Contract Start derived from driving_license_date, has an average of ~21 years.
While it is highly relevant for assessing risk, the data contains negative values and extreme outliers that indicate quality issues.
These anomalies should be treated carefully before using the feature in models.
"""

# DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='feature_2_insights', value=feature_2_insights)
```

feature_2_insights

Driving Experience (years) at Contract Start derived from driving_license_date, has an average of ~21 years, with most customers between 11 and 30 years. While it is highly relevant for assessing risk, the data contains negative values and extreme outliers that indicate quality issues. These anomalies should be treated carefully before using the feature in models.

C.4 Explore Feature of Interest seniority

You can add more cells in this section

```
In [32]: feature_name      = "seniority"          # Years associated with the insurer
ref_date_col       = "contract_start_date"    # Reference date column (parsed for cohort views)
target_name        = "net_premium_amount"       # Target variable
print_feature_name = "Customer Seniority (years) with the Insurer"
```

```
In [33]: # Prepare data
sen = pd.to_numeric(training_df[feature_name], errors="coerce")
ref_dt = pd.to_datetime(training_df[ref_date_col], errors="coerce") # kept for potential cohort analysis

# Descriptive statistics
print(f"Feature: {feature_name} -> {print_feature_name}")
print(sen.describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]))

# Quick sanity notes (no hard filtering)
if (sen < 0).any(): print("Note: Negative seniority values exist (data quality).")
if (sen > 50).any(): print("Note: Very large seniority values (>50) exist (check plausibility)")

# Distribution plots
```

```

valid = sen.dropna()
if len(valid) > 0:
    plt.figure(figsize=(10,5))
    plt.hist(valid, bins=30)
    plt.title(print_feature_name)
    plt.xlabel("Years"); plt.ylabel("Count")
    plt.tight_layout(); plt.show()

    plt.figure(figsize=(8,3))
    plt.boxplot(valid, vert=False)
    plt.title(f"Boxplot: {print_feature_name}")
    plt.xlabel("Years")
    plt.tight_layout(); plt.show()

# Relationship with target (scatter + grouped averages)
df_plot = pd.DataFrame({ "Seniority": sen, target_name: training_df[target_name] }).dropna()

# Scatter
plt.figure(figsize=(8,5))
plt.scatter(df_plot["Seniority"], df_plot[target_name], alpha=0.3)
plt.title(f"{print_feature_name} vs {target_name}")
plt.xlabel("Seniority (years)"); plt.ylabel(target_name)
plt.tight_layout(); plt.show()

# Grouped averages (bins can be adjusted)
# Uses broader bins to stabilize means across skewed counts
lo, hi = np.nanpercentile(df_plot["Seniority"], [1, 99])
bins = [0, 1, 2, 3, 5, 10, 20, max(40, hi)]
df_plot["SenGroup"] = pd.cut(df_plot["Seniority"], bins=bins, right=False, include_lowest=True)
mean_prem = df_plot.groupby("SenGroup")[target_name].mean()

plt.figure(figsize=(9,4))
plt.bar(mean_prem.index.astype(str), mean_prem.values)
plt.title(f"Mean {target_name} by {print_feature_name} Group")
plt.ylabel(f"Mean {target_name}"); plt.xlabel("Seniority Group (years)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout(); plt.show()


```

Feature: seniority -> Customer Seniority (years) with the Insurer


count	32136.000000
mean	5.896596
std	5.909002
min	1.000000
1%	1.000000
5%	1.000000
25%	2.000000
50%	4.000000
75%	7.000000
95%	19.000000
99%	29.000000
max	40.000000

Name: seniority, dtype: float64


Customer Seniority (years) with the Insurer




Boxplot: Customer Seniority (years) with the Insurer



Customer Seniority (years) with the Insurer vs net_premium_amount





```
In [34]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues, etc.

feature_3_insights = """
Seniority (years) has an average of about 6 years, with most customers between 2 and 7 years. The distribution is right-skewed, showing many recent customers and fewer long-term ones, with a maximum of 40 years. This feature is relevant as longer seniority may indicate loyalty and lower churn risk. However, it may overlap with policy history variables, leading to multicollinearity. Extreme values (30+ years) are rare and should be checked for plausibility before modeling.
"""

```

```
In [35]: # DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='feature_3_insights', value=feature_3_insights)
```

feature_3_insights

Seniority (years) has an average of about 6 years, with most customers between 2 and 7 years. The distribution is right-skewed, showing many recent customers and fewer long-term ones, with a maximum of 40 years. This feature is relevant as longer seniority may indicate loyalty and lower churn risk. However, it may overlap with policy history variables, leading to multicollinearity. Extreme values (30+ years) are rare and should be checked for plausibility before modeling.

C.4 Explore Feature of Interest policy_type

You can add more cells related to other features in this section

```
In [36]: feature_name      = "policy_type"           # Categorical feature (e.g., 1/2/3/4)
ref_date_col       = "contract_start_date"    # Reference date column
target_name        = "net_premium_amount"       # Target variable
print_feature_name = "Policy Type"
```

```
In [37]: # Prepare data
s = training_df[feature_name]
ref_dt = pd.to_datetime(training_df[ref_date_col], errors="coerce") # parsed for completeness
t = training_df[target_name]

# Label NaN for plotting
s_plot = s.astype(object).where(s.notna(), "Missing")

# Basic stats
```

```

n = len(s); miss = s.isna().sum(); uniq = s.unique(dropna=True)
print(f"Feature: {feature_name} -> {print_feature_name}")
print(f"Non-null: {n - miss}/{n} (missing: {miss}, {miss/n:.2%})")
print(f"Unique categories (non-null): {uniq}")
print("\nCounts (incl. Missing):")
print(s_plot.value_counts(dropna=False))

# Count bar plot
counts = s_plot.value_counts(dropna=False)
plt.figure(figsize=(8,4))
plt.bar(counts.index.astype(str), counts.values)
plt.title(f"{print_feature_name} - Count"); plt.xlabel(print_feature_name); plt.ylabel("Count")
plt.tight_layout()
plt.show()

# Mean target by category (ordered by mean)
df_plot = pd.DataFrame({print_feature_name: s_plot, target_name: t}).dropna(subset=[print_feature_name])
mean_target = df_plot.groupby(print_feature_name)[target_name].mean().sort_values(ascending=False)

plt.figure(figsize=(8,4))
plt.bar(mean_target.index.astype(str), mean_target.values)
plt.title(f"Mean {target_name} by {print_feature_name}")
plt.xlabel(print_feature_name); plt.ylabel(f"Mean {target_name}")
plt.tight_layout()
plt.show()

# Target distribution by category (boxplots)
# Build a list of arrays aligned with categories (use top-k if too many)
cats = counts.index.tolist()
data_by_cat = [df_plot.loc[df_plot[print_feature_name] == c, target_name].values for c in cats]

plt.figure(figsize=(9,4))
plt.boxplot(data_by_cat, labels=[str(c) for c in cats], vert=True, showfliers=False)
plt.title(f"{target_name} Distribution by {print_feature_name} (no outliers)")
plt.xlabel(print_feature_name); plt.ylabel(target_name)
plt.tight_layout()
plt.show()


```

Feature: policy_type -> Policy Type
 Non-null: 32136/32136 (missing: 0, 0.00%)
 Unique categories (non-null): 3

Counts (incl. Missing):

policy_type	
3	27348
2	4428
1	360

Name: count, dtype: int64



```
In [38]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues, etc.

feature_4_insights = """
Policy Type is a categorical variable with 3 categories present in the dataset:
type 3 (passenger cars) dominates with ~85% of records, followed by type 2 (vans) and a very low count for type 1 (motorcycles).
Boxplots show that average premiums differ slightly across policy types, with type 2 policies having the highest premiums on average.
This suggests the feature has predictive value for premiums.
"""

print(feature_4_insights)
```

The sample size of type 1 is small. It may require grouping or regularization in modeling.

""

In [39]: # DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='feature_4_insights', value=feature_4_insights)

feature_4_insights

Policy Type is a categorical variable with 3 categories present in the dataset: type 3 (passenger cars) dominates with ~85% of records, followed by type 2 (vans) and a very small group of type 1 (motorbikes). Boxplots show that average premiums differ slightly across policy types, with type 2 policies generally priced higher and type 1 lower. This suggests the feature has predictive value for premiums. The sample size of type 1 is small. It may require grouping or regularization in modeling.

C.5 Explore Feature of Interest `total_claims_number_in_history`

You can add more cells related to other features in this section

In [40]:

```
feature_name      = "total_claims_number_in_history"    # Historical claims count
ref_date_col      = "contract_start_date"                # Reference date (parsed for completeness)
target_name        = "net_premium_amount"                 # Target variable
print_feature_name = "Historical Claims Count"
```

In [41]:

```
# Prepare data
x = pd.to_numeric(training_df[feature_name], errors="coerce")
ref_dt = pd.to_datetime(training_df[ref_date_col], errors="coerce") # not used further, but included
y = training_df[target_name]

# Descriptive stats
print(f"Feature: {feature_name} -> {print_feature_name}")
print(x.describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]))

# Distribution (hist + boxplot)
valid = x.dropna()
if len(valid) > 0:
    plt.figure(figsize=(10,5))
    plt.hist(valid, bins=30)
    plt.title(f"{print_feature_name} - Distribution")
    plt.xlabel(print_feature_name); plt.ylabel("Count")
    plt.tight_layout(); plt.show()

    plt.figure(figsize=(8,3))
    plt.boxplot(valid, vert=False, showfliers=True)
    plt.title(f"Boxplot: {print_feature_name}")
    plt.xlabel(print_feature_name)
    plt.tight_layout(); plt.show()

# Relationship with target (scatter + grouped averages)
df_plot = pd.DataFrame({print_feature_name: x, target_name: y}).dropna()

# Scatter (alpha to mitigate overlap)
plt.figure(figsize=(8,5))
plt.scatter(df_plot[print_feature_name], df_plot[target_name], alpha=0.3)
plt.title(f"{print_feature_name} vs {target_name}")
plt.xlabel(print_feature_name); plt.ylabel(target_name)
plt.tight_layout(); plt.show()

# Grouped mean premium by claims buckets (bins can be adjusted)
```

```

# Buckets: 0, 1, 2-3, 4-6, 7-10, >10
bins = [-0.5, 0.5, 1.5, 3.5, 6.5, 10.5, df_plot[print_feature_name].max()+0.5]
labels = ["0", "1", "2-3", "4-6", "7-10", ">10"]
df_plot["ClaimsBucket"] = pd.cut(df_plot[print_feature_name], bins=bins, labels=labels)

mean_prem = df_plot.groupby("ClaimsBucket")[target_name].mean()
plt.figure(figsize=(8,4))
plt.bar(mean_prem.index.astype(str), mean_prem.values)
plt.title(f"Mean {target_name} by {print_feature_name} Bucket")
plt.ylabel(f"Mean {target_name}"); plt.xlabel("Claims Bucket")
plt.tight_layout(); plt.show()


```

Feature: total_claims_number_in_history -> Historical Claims Count


count	32136.000000
mean	2.647498
std	3.787398
min	0.000000
1%	0.000000
5%	0.000000
25%	0.000000
50%	1.000000
75%	4.000000
95%	10.000000
99%	17.000000
max	52.000000

Name: total_claims_number_in_history, dtype: float64


Historical Claims Count - Distribution




Boxplot: Historical Claims Count



Historical Claims Count vs net_premium_amount



Mean net_premium_amount by Historical Claims Count Bucket



```
In [42]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues,  
feature_5_insights = """  
Historical Claims Count(total_claims_number_in_history) has a mean of about 2.65 claims, with  
The distribution is highly right-skewed, with most values clustered near zero and a few extre  
This feature is important for capturing past risk behavior, as higher claim counts usually inc  
Issues: There exist some extreme values(up to 52 claims), it should be capped or grouped into  
Overall, the feature is predictive but requires handling of skewness and outliers.  
"""
```

```
In [43]: # DO NOT MODIFY THE CODE IN THIS CELL  
print_tile(size="h3", key='feature_5_insights', value=feature_5_insights)
```

Historical Claims Count(total claims number in history) has a mean of about 2.65 claims, with half of customers reporting one or fewer claims. The distribution is highly right-skewed, with most values clustered near zero and a few extreme outliers up to 52 claims. This feature is important for capturing past risk behavior, as higher claim counts usually indicate higher risk and premiums. Issues: There exist some extreme values(up to 52 claims), it should be capped or grouped into buckets. Overall, the feature is predictive but requires handling of skewness and outliers.

C.6 Explore Feature of Interest vehicle_value

You can add more cells related to other features in this section

```
In [44]: feature_name      = "vehicle_value"          # Vehicle market value
ref_date_col       = "contract_start_date"    # Reference date (parsed for completeness)
target_name        = "net_premium_amount"       # Target variable
print_feature_name = "Vehicle Value"
```

```
In [45]: # Prepare data
x = pd.to_numeric(training_df[feature_name], errors="coerce")
ref_dt = pd.to_datetime(training_df[ref_date_col], errors="coerce")    # parsed to satisfy ref
y = training_df[target_name]

# Descriptive stats
print(f"Feature: {feature_name} -> {print_feature_name}")
print(x.describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]))

# Quick sanity notes (no hard filtering)
if (x <= 0).any():
    print("Note: non-positive values detected; check data quality or consider imputation.")
if x.isna().any():
    print(f"Note: {x.isna().sum()} missing values; consider median/grouped imputation.")

# Distribution plots
valid = x.dropna()
if len(valid) > 0:
    # Histogram
    plt.figure(figsize=(10,5))
    plt.hist(valid, bins=30)
    plt.title(f"{print_feature_name} - Distribution")
    plt.xlabel(print_feature_name); plt.ylabel("Count")
    plt.tight_layout(); plt.show()

    # Boxplot (shows outliers)
    plt.figure(figsize=(8,3))
    plt.boxplot(valid, vert=False, showfliers=True)
    plt.title(f"Boxplot: {print_feature_name}")
    plt.xlabel(print_feature_name)
    plt.tight_layout(); plt.show()

# Log1p view for right-skew
plt.figure(figsize=(10,5))
plt.hist(np.log1p(valid), bins=30)
plt.title(f"log1p({print_feature_name}) - Distribution")
plt.xlabel(f"log1p({print_feature_name})"); plt.ylabel("Count")
plt.tight_layout(); plt.show()

# Relationship with target (scatter + bucketed means)
```

```

df_plot = pd.DataFrame({print_feature_name: x, target_name: y}).dropna()

# Scatter
plt.figure(figsize=(8,5))
plt.scatter(df_plot[print_feature_name], df_plot[target_name], alpha=0.3)
plt.title(f"{print_feature_name} vs {target_name}")
plt.xlabel(print_feature_name); plt.ylabel(target_name)
plt.tight_layout(); plt.show()

# Bucketed averages using quantile-based bins (robust to skew)
qs = np.unique(np.quantile(df_plot[print_feature_name], [0, .1, .25, .5, .75, .9, 1.0]))
# ensure strictly increasing edges
bins = np.unique(qs)
if len(np.unique(bins)) < 3:
    # fallback fixed bins if quantiles collapse (rare)
    bins = [df_plot[print_feature_name].min()-1, 10000, 15000, 20000, 30000, df_plot[print_feature_name].max()]
df_plot["ValueBucket"] = pd.cut(df_plot[print_feature_name], bins=bins, include_lowest=True)
mean_prem = df_plot.groupby("ValueBucket")[target_name].mean()


plt.figure(figsize=(9,4))
plt.bar(mean_prem.index.astype(str), mean_prem.values)
plt.title(f"Mean {target_name} by {print_feature_name} Bucket")
plt.ylabel(f"Mean {target_name}"); plt.xlabel(f"{print_feature_name} Bucket")
plt.xticks(rotation=45, ha="right")
plt.tight_layout(); plt.show()

```


Feature: vehicle_value -> Vehicle Value

count	32136.000000
mean	18972.391870
std	7176.732691
min	270.460000
1%	7068.000000
5%	9900.000000
25%	14093.730000
50%	18020.000000
75%	22466.050000
95%	31200.000000
99%	42746.250000
max	110000.000000


Name: vehicle_value, dtype: float64




Boxplot: Vehicle Value




log1p(Vehicle Value) - Distribution



Vehicle Value vs net_premium_amount





```
In [46]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues, etc.

feature_6_insights = """
Vehicle Value has an average of about 19,000 with most values ranging between 14,000 and 22,500. The distribution is strongly right-skewed, which was confirmed by the log transformation produced. This feature is important because higher vehicle value generally corresponds to higher replacement cost and risk, making it highly predictive for premium pricing. Issues: There exist some extreme values, it should be addressed through log transformation or outlier capping.
"""

# DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='feature_6_insights', value=feature_6_insights)
```

```
feature_6_insights
```

Vehicle Value has an average of about 19,000 with most values ranging between 14,000 and 22,500, but extreme outliers reach up to 110,000. The distribution is strongly right-skewed, which was confirmed by the log transformation producing a near-normal shape. This feature is important because higher vehicle value generally corresponds to higher replacement cost and risk, making it highly predictive for premium pricing. Issues: There exist some extreme values, it should be addressed through log transformation or outlier capping.

C.7 Explore Feature of Interest matriculation_year

You can add more cells related to other features in this section

```
In [48]: feature_name      = "matriculation_year"          # Vehicle registration year (YYYY)
ref_date_col     = "contract_start_date"           # Reference date (contract start)
target_name       = "net_premium_amount"            # Target variable
print_feature_name = "Car Age (years) at Contract Start"
```

```
In [49]: # Prepare data
year = pd.to_numeric(training_df[feature_name], errors="coerce")
ref = pd.to_datetime(training_df[ref_date_col], errors="coerce")

# Compute car age at contract start (allow NaNs)
car_age = (ref.dt.year - year).where(year.notna() & ref.notna(), np.nan)

# Descriptive stats
print(f"Feature: {feature_name} -> {print_feature_name}")
```

```

print(car_age.describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]))
if (car_age < 0).any(): print("Note: negative car ages exist (check registration vs contract start date)")

# Distribution plots
valid = car_age.dropna()
plt.figure(figsize=(10,5))
plt.hist(valid, bins=30)
plt.title(print_feature_name)
plt.xlabel("Years"); plt.ylabel("Count")
plt.tight_layout(); plt.show()

plt.figure(figsize=(8,3))
plt.boxplot(valid, vert=False)
plt.title(f"Boxplot: {print_feature_name}")
plt.xlabel("Years")
plt.tight_layout(); plt.show()

# Relationship with target (scatter + bucketed means)
df_plot = pd.DataFrame({print_feature_name: car_age, target_name: training_df[target_name]}).reset_index()

# Scatter
plt.figure(figsize=(8,5))
plt.scatter(df_plot[print_feature_name], df_plot[target_name], alpha=0.3)
plt.title(f"{print_feature_name} vs {target_name}")
plt.xlabel(print_feature_name); plt.ylabel(target_name)
plt.tight_layout(); plt.show()



# Quantile buckets (robust to skew)
qs = np.unique(np.quantile(df_plot[print_feature_name], [0, .1, .25, .5, .75, .9, 1.0]))
if len(qs) < 3: # fallback if degenerate
    qs = [df_plot[print_feature_name].min()-1, 3, 5, 10, 15, 20, df_plot[print_feature_name].max()]
df_plot["AgeBucket"] = pd.cut(df_plot[print_feature_name], bins=qs, include_lowest=True)
mean_prem = df_plot.groupby("AgeBucket")[target_name].mean()


plt.figure(figsize=(9,4))
plt.bar(mean_prem.index.astype(str), mean_prem.values)
plt.title(f"Mean {target_name} by {print_feature_name} Bucket")
plt.ylabel(f"Mean {target_name}"); plt.xlabel("Car Age Bucket (years)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout(); plt.show()

```

Feature: matriculation_year -> Car Age (years) at Contract Start

count	32136.000000
mean	15.793316
std	5.971988
min	6.000000
1%	6.000000
5%	6.000000
25%	11.000000
50%	16.000000
75%	20.000000
95%	25.000000
99%	30.000000
max	70.000000
dtype:	float64





```
In [50]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues, etc.

feature_7_insights = """
Car Age (years) at Contract Start has a mean of about 16 years, with most vehicles between 11 and 20 years old. The distribution is moderately right-skewed, with a few very old vehicles up to 70 years that may be data anomalies or rare cases. A limitation is that car age may correlate with vehicle value, multicollinearity should be considered. Extreme values should be reviewed, and the feature may be more useful in bucketed form rather than as a raw continuous variable.
"""

# DO NOT MODIFY THE CODE IN THIS CELL
print_tile(size="h3", key='feature_7_insights', value=feature_7_insights)
```

feature_7_insights

Car Age (years) at Contract Start has a mean of about 16 years, with most vehicles between 11 and 20 years old. The distribution is moderately right-skewed, with a few very old vehicles up to 70 years that may be data anomalies or rare cases. A limitation is that car age may correlate with vehicle value, multicollinearity should be considered. Extreme values should be reviewed, and the feature may be more useful in bucketed form rather than as a raw continuous variable.

C.8 Explore Feature of Interest `vehicle_fuel_type`

You can add more cells related to other features in this section

```
In [52]: feature_name      = "vehicle_fuel_type"      # Fuel type (e.g., Petrol/Diesel)
ref_date_col       = "contract_start_date"    # Reference date (parsed for completeness)
target_name        = "net_premium_amount"       # Target variable
print_feature_name = "Vehicle Fuel Type"
```

```
In [53]: # Prepare data
s = training_df[feature_name].astype(str).where(training_df[feature_name].notna(), "Missing")
ref_dt = pd.to_datetime(training_df[ref_date_col], errors="coerce") # parsed but unused
y = training_df[target_name]

# Basic stats
print(f"Feature: {feature_name} -> {print_feature_name}")
print("Counts (incl. Missing):")
print(s.value_counts())
```

```

# Count bar plot
counts = s.value_counts()
plt.figure(figsize=(6,4))
plt.bar(counts.index.astype(str), counts.values)
plt.title(f"{print_feature_name} - Count")
plt.xlabel(print_feature_name); plt.ylabel("Count")
plt.tight_layout(); plt.show()

# Mean premium by fuel type
df_plot = pd.DataFrame({print_feature_name: s, target_name: y}).dropna()
mean_target = df_plot.groupby(print_feature_name)[target_name].mean().sort_values(ascending=False)

plt.figure(figsize=(6,4))
plt.bar(mean_target.index.astype(str), mean_target.values)
plt.title(f"Mean {target_name} by {print_feature_name}")
plt.xlabel(print_feature_name); plt.ylabel(f"Mean {target_name}")
plt.tight_layout(); plt.show()

# Premium distribution by category (boxplot)
groups = [df_plot.loc[df_plot[print_feature_name] == cat, target_name].values for cat in mean_target.index]

```

Feature: vehicle_fuel_type -> Vehicle Fuel Type

Counts (incl. Missing):

vehicle_fuel_type


D 21450



P 10608

Missing 78

Name: count, dtype: int64

Vehicle Fuel Type - Count





```
In [54]: # provide a detailed analysis on the selected feature, its distribution, limitations, issues,  
feature_8_insights = """  
The feature Vehicle Fuel Type is dominated by Diesel (~67%) and Petrol (~33%), with very few vehicles having CNG.  
The distribution of premiums shows Diesel vehicles having slightly higher average premiums than Petrol vehicles.  
The Missing group is lower but very small in size.  
Issues: there exist small number of missing values, which may be treated as a separate category.  
"""
```

```
In [55]: # DO NOT MODIFY THE CODE IN THIS CELL  
print_tile(size="h3", key='feature_8_insights', value=feature_8_insights)
```

The feature Vehicle Fuel Type is dominated by Diesel ($\approx 67\%$) and Petrol ($\approx 33\%$), with very few missing cases. The distribution of premiums shows Diesel vehicles having slightly higher average premiums than Petrol. The Missing group is lower but very small in size. Issues: there exist small number of missing values, which may be treated as a separate category or imputed.